

# Algorithm Engineering Lab Assignment 2

Brian Zahoransky (brian.zahoransky@uni-jena.de)

January 15, 2021

## 1. What causes false sharing?

---

False sharing may appear if two processes are going to write to the same cache line at once. For further explanation consider the following example. A program should show that floating point operations are not associative by summing up a huge number of floating point values. First, an array containing the values which should be summed is initialized. Second, the variables sum1 and sum2 are initialized with zero. Both steps happen, while the program runs serially. Assume that both variables, sum1 and sum2, are stored in the same cache line. In sum1 will be stored the result summing up all values in right order, sum2 the algorithm sums the float array in inverse order. To speed up the program, each summation runs on a own thread on a different processor core. Actually, there is no speed up, more likely is a slow down because of false sharing. After every update of sum1 or sum2 the cache line stored in the cache of the other processor becomes invalid. Thus, the updated cache line has to be send to the other processor before it can proceed. Since a huge array of floats is assumed, most time the processors will wait for a valid cache line.

## 2. How do mutual exclusion constructs prevent race conditions?

---

Race conditions occur if multiple threads try to access one resource at once. In the context of storage, a race condition occurs if multiple threads write to the same memory address. For example, multiple threads performing an operation on a global one. First, thread A loads the value of the global variable, performs its operation. Now, thread A stops for some reason, maybe because the system decides to give B computational time but other scenarios are plausible as well. However, B also loads an old copy of the global variable, performs its operation, and writes the update back into storage. Afterwards, thread A gets computational time again. It stores its result of the global variable and overwrites the changes B has made.

To avoid such behavior mutual exclusions can be used. They allow only one thread to access a critical resource. Consider the example above. Thread A wants to update the global variable. Therefore, it receives the access right by locking the mutex. Note that locking and freeing the mutex is a atomic operation, thus it cannot be interrupted. If the computational time of A stops now, no other thread is able to change the value of the global variable. Even if B proceeds it stops before updating the global variable. However, sometime thread A gets computational time again. It can perform its operation, write back the result, and free the mutex. Now, B will be able to lock the mutex, gets the new version of the global variable, performs its operation, and stores the result.

## 3. Explain the differences between static and dynamic schedules in OpenMP.

---

**static:** At compile time is the schedule fix. Thus, each chunk is assigned to a specific thread after the compile procedure. It is applied when each chunk holds nearly the same amount of work. The software engineer may consider using a static schedule if the workload differs a lot and can be estimated well.

**dynamic:** The chunks are assigned to specific threads at run time. It is applied when the workload for the chunks are unknown or the chunks cannot be divided similarly using a static schedule. However, performing the scheduling task at run time causes performance penalties. This should be considered by the choice of the used schedule.

---

**4. What can we do if we have found a solution while running a parallel for loop in OpenMP, but still have many iterations left?**

---

One solution can be realized by using two atomic variables. One is a Boolean initialized with false. It will be changed to true once a solution is found. The first line of the parallel for loop is an if-clause. Its condition is the atomic Boolean. Once the Boolean is set to true in each iteration happens nothing but the evaluation of the if-clause and the following continue instruction. While the Boolean is false the search proceeds, which means executing the code the ending of the if-clause. A if-clause later on checks if a solutions has be found. When that is the case, the atomic Boolean is set to true and the result is written into the atomic solution-variable.

---

**5. Explain in your own words how `std::atomic::compare_exchange_weak` work.**

---

`value.compare_exchange_weak(expected_value, desired_value);`

Given the code line above, the atomic variable *value* should become overwritten. Therefore, *value* is compared with *expected\_value*. If *value* and *expected\_value* are equal, the *desired\_value* is written into *value* and the function returns true. If *value* and *expected\_value* are not equal, the *desired\_value* is written into *expected\_value* and the function returns false. The second scenario may occur when another thread has changed the *value* after the *expected\_value* was set and before the `compare_exchange_weak` function has been invoked.