

Algorithm Engineering Lab Assignment 4

Brian Zahoransky (brian.zahoransky@uni-jena.de)

February 22, 2021

1. Explain how divide and conquer algorithms can be parallelized with tasks in OpenMP.

The goal of divide-and-conquer algorithms is to generate an ordered list out of a randomly sorted one. The algorithm can be described through a recursive function. This function receives a not ordered list of elements. If the length of this list is one, it is ordered and is returned. Otherwise, the list of elements is divided into parts. Now, the function calls itself for each sublist. After this phase, the function holds ordered sublists. The sublists become merged to generate a fully ordered list. The result is returned, and the algorithm can be parallelised with the help of tasks. A new task starts if the function calls itself as long as the list length is above a set value. This mechanism ensures that the overhead for starting new tasks does not exceed the work for sorting a list.

2. Describe some ways to speed up merge sort.

Alternative sorting algorithms for small sub-lists: For example, insertion-sort is faster than merge-sort for small lists. Due to this, it was used in the lecture for lists with less than 32 elements.

Storing in stack instead of heap: The storage for stacks is faster and limited. Therefore, lists are stored in stacks when their length was below 8192 in the lecture.

Multiple Task algorithms: An additional function is used to enable a multiple task version of merge-sort. This function opens a parallel region and invokes merge-sort. That and the omp single instruction within the merge-sort function avoid that each piece of work is processed from multiple tasks. For further descriptions of multi-threaded merge-sort, have a look at exercise 1 and exercise 3.

In-place algorithms: The approach is working with two lists of the same size. One list holds the elements to be sorted; the other one is a buffer. Both lists are declared before the first call of the merge-function. Thus, there is no need for creating new lists during the execution of merge-sort.

3. What is the idea behind multi-threaded merging?

A recursive function gets two sorted sublists A and B, merges them into a sorted list, and returns it. Through a lookup on the longer sub-list, its median M is determined. Now, the two lists become divided by M in A_smaller_M, A_bigger_M, B_smaller_M, and B_bigger_M. Afterwards, the function calls itself, once with A_smaller_M and B_smaller_M, and once with A_bigger_M and B_bigger_M. The results are two sorted lists M1 and M2. M1 contains only elements smaller than M, and M2 only elements bigger than M. Concatenating M1, M, and M2 generates a fully sorted list. As described in exercise 1, tasks can be started each time the function calls itself.

4. Read What every systems programmer should know about concurrency (<https://assets.bitbashing.io/papers/concurrency-primer.pdf>). Discuss two things you find particularly interesting.

In the tenth chapter, memory orderings are discussed. A programmer can give each atomic operation an optional order-clause. The paper also describes use cases. However, the following only explains the procedures with the help of pseudo-code:

```
doWork1()
barrier1()
atomicOperation()
barrier2()
doWork2()
```

memory_order_seq_cst: It stands for sequentially consistent order and is set per default. Above, a pseudo-code example is shown. To improve performance, the compiler and processor are only allowed to reorder the operations within the blocks `doWork1()` and `doWork2()`. However, they are not allowed to reorder work through the barriers. First, `doWork1()` has to be done, then the atomic operation, and last `doWork2()`.

memory_order_release: This clause removes the `barrier1()`. It allows the compiler and processor to reorder `doWork1()` and the atomic operation.

memory_order_acquire: It is usually the counterpart to the release clause in a second thread. A further explanation can be found in the paper. Anyway, this clause keeps `barrier1()` and removes `barrier2()`. Thus, `doWork2()` and the atomic operation may be shuffled by a compiler or processor to increase the performance.

memory_order_acq_rel: Both barriers are kept. The difference compared to the sequential consistent case is described in the paper as follows, "acquire-release provides order relative to the variable being load-acquired and store-released, whereas sequentially consistent operation provides some global order across the entire program."

memory_order_consume: It is the alternative counterpart for release and works similar to acquire. The difference between consuming and acquire is the interpretation of the kept barrier. The consume clause let operations pass the barrier as long as they are not dependent on the atomic variable which was affected by the `atomicOperation()`.

memory_order_relaxed: That removes both barriers. One can use it if the operation has only to be atomic, but the order does not matter. Thus, the compiler and processor are allowed to reorder `doWork1()`, `atomicOperation()`, and `doWork2()` to speed up the program.

Another particular interesting topic is the use of volatile keyword (chapter 13). Besides, it has already been discussed in the lecture; the paper explains it from a different point of view. It cannot be applied to build concurrent generally. Volatile implies two guarantees: The compiler will not remove loads and stores that seem unnecessary, and the compiler will not reorder volatile reads / writes with respect to other ones.